



MAGIC CODE INSTITUTE



## Lecture 6: Modeling Libraries in Python



# MAGIC CODE INSTITUTE

## Section 6.1 | Interfacing Between pandas and Model Code

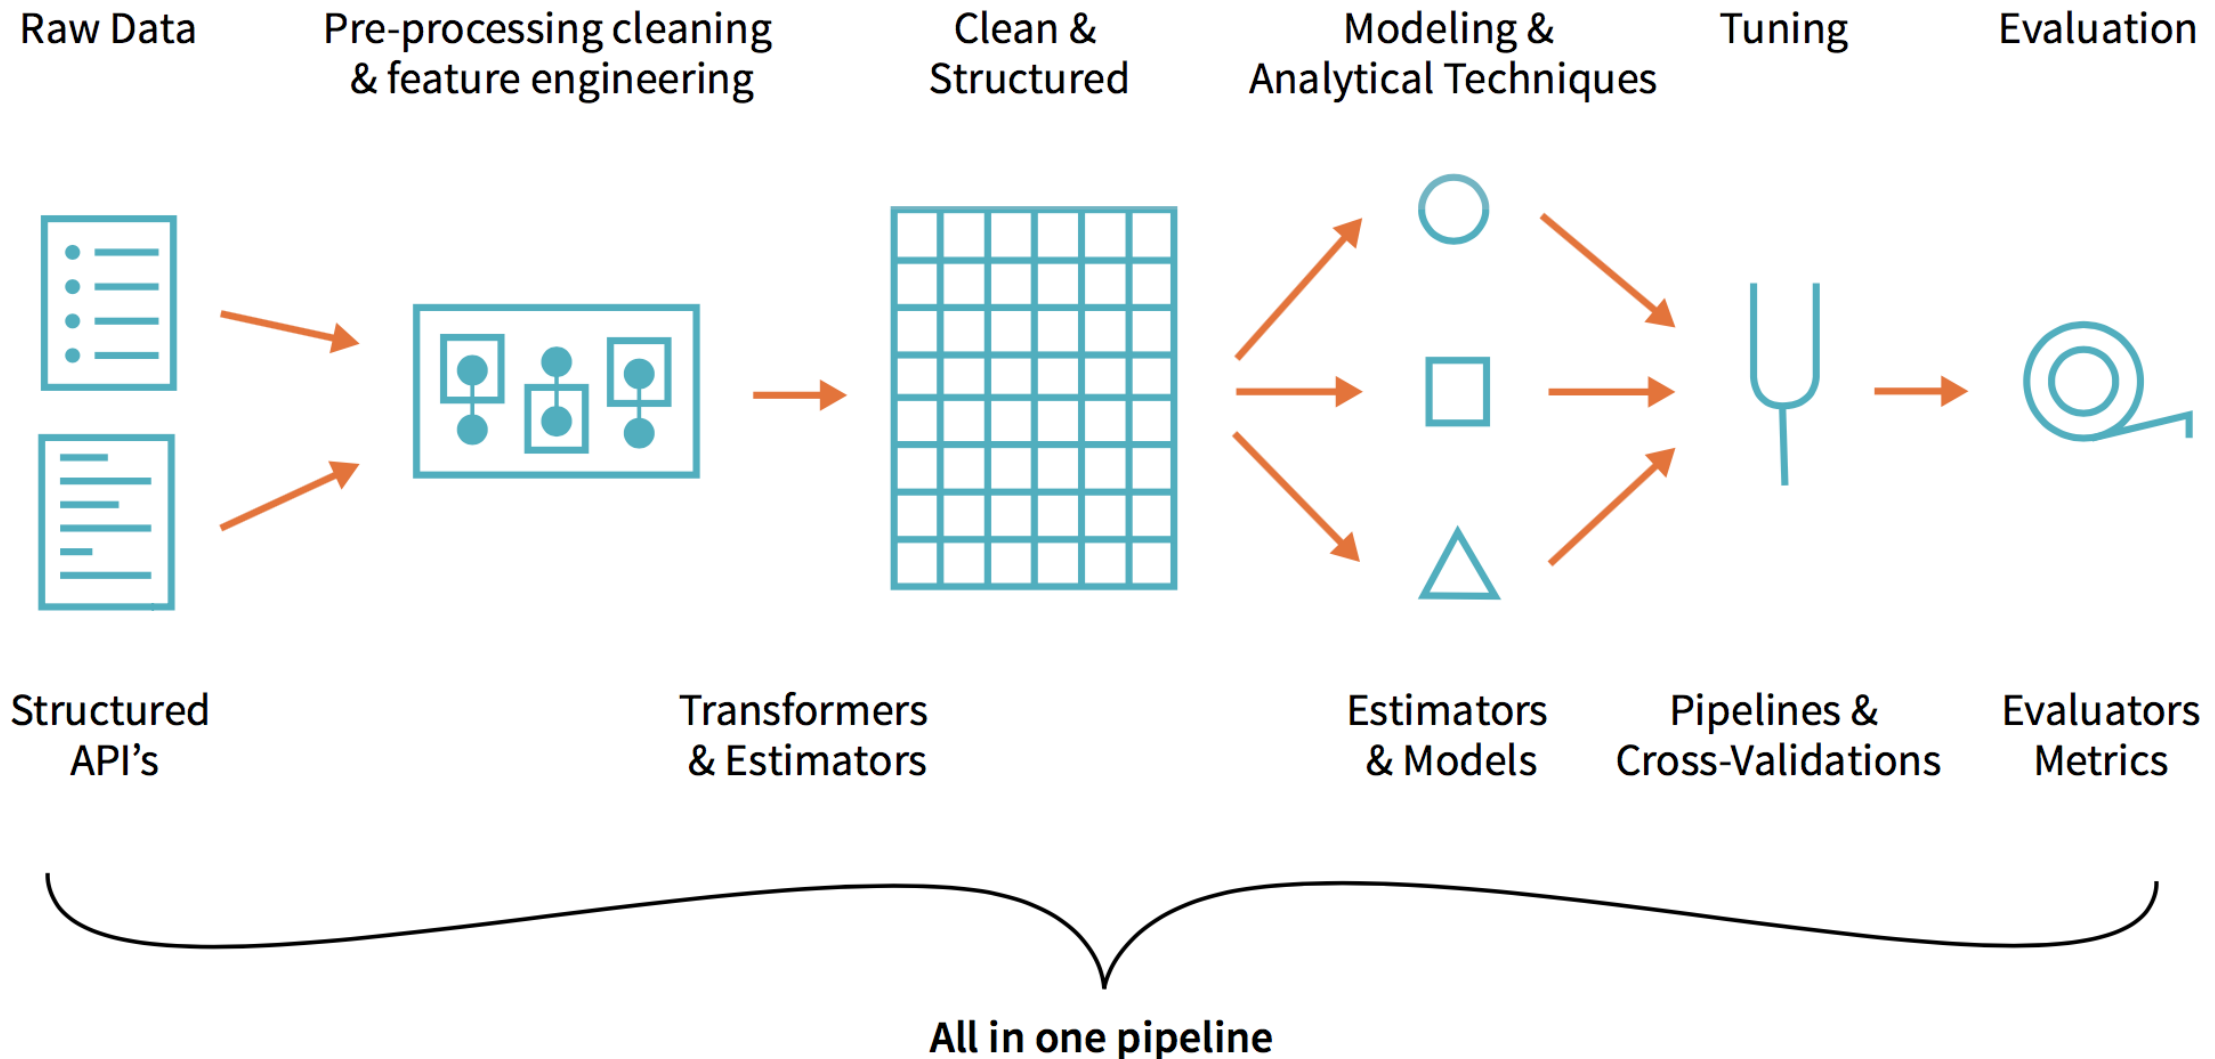
■

SCRIPT





## ➤ Progress of a data science project:





# FEATURE ENGINEERING

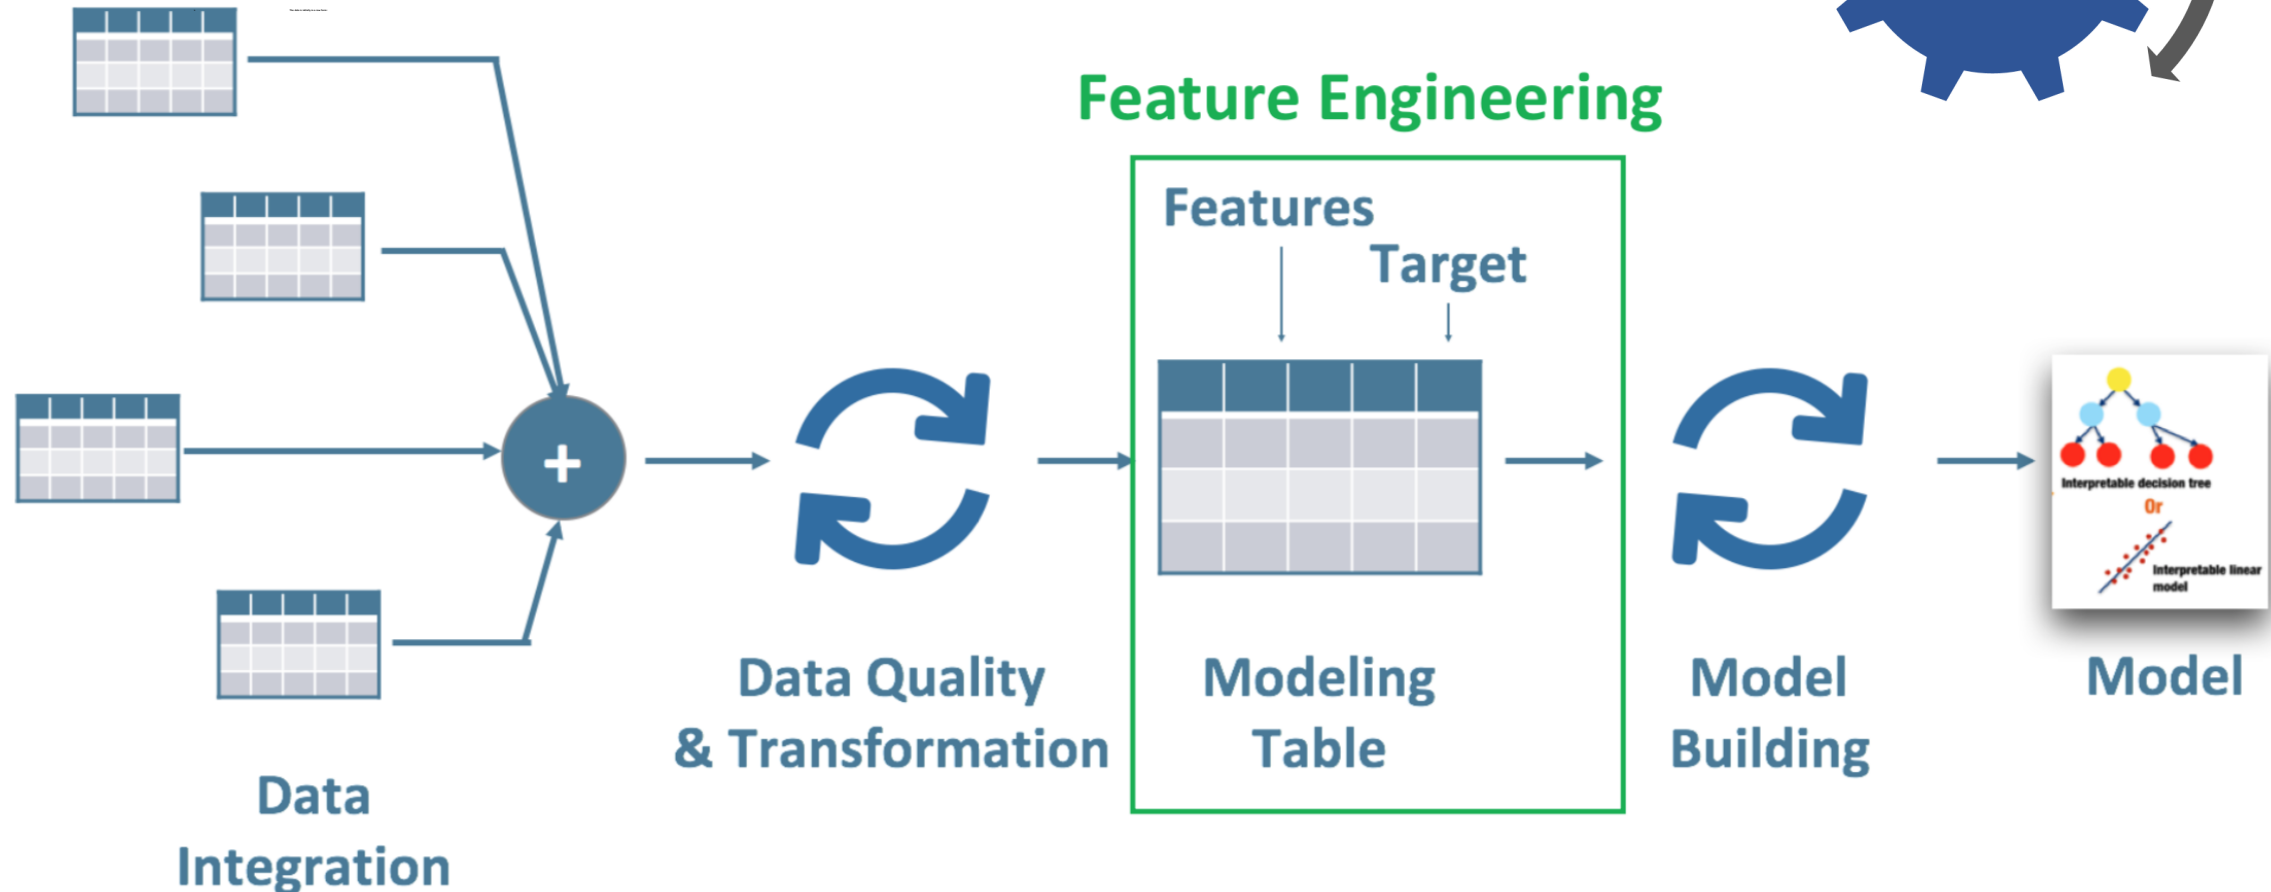
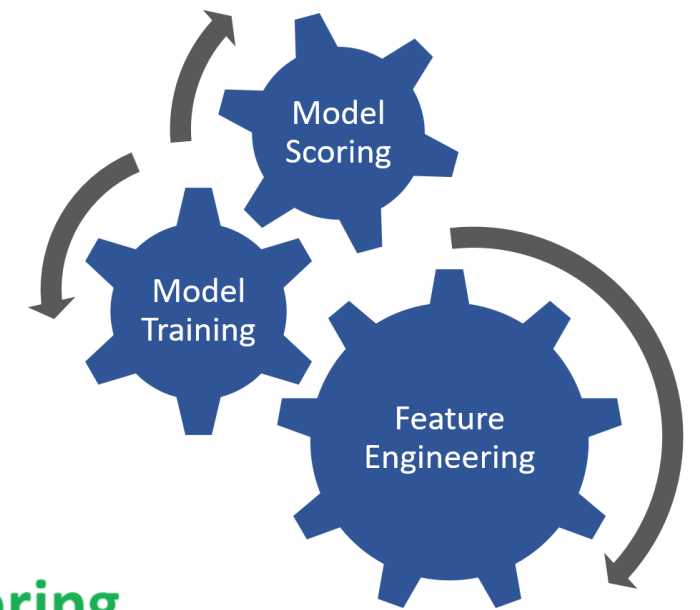






## Why Feature Engineering is important:

- Data is initially in a raw form and not ready for analyzing / modeling purposes
- Feature engineering is the process of using domain knowledge of the data to create (hopefully) features
- Account for 80% - 90% of a data science project



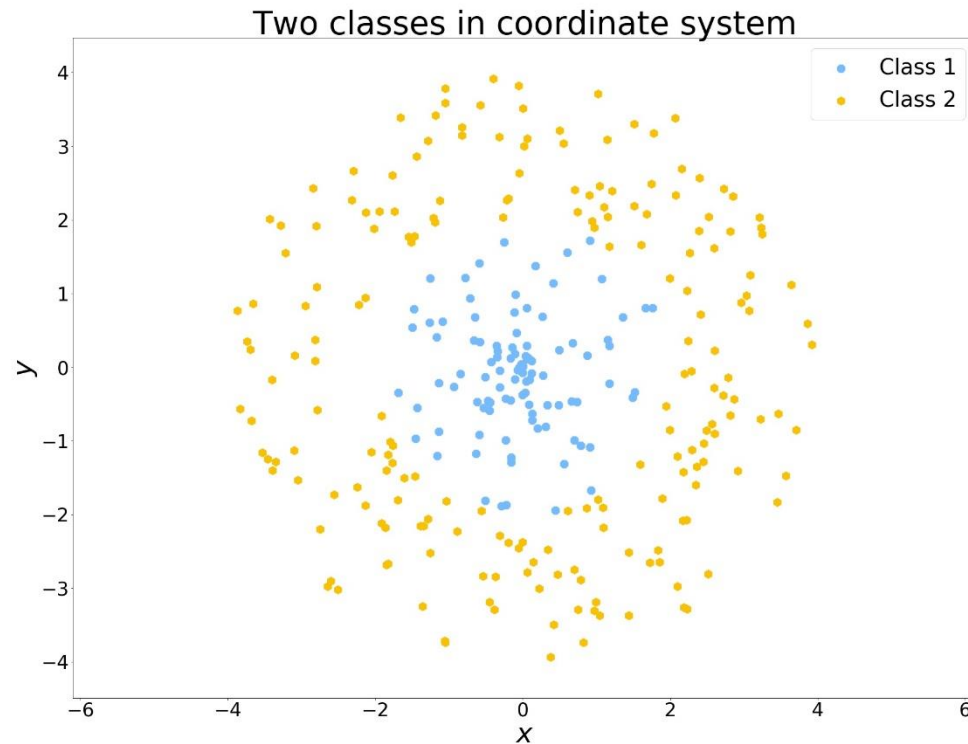


### What is Feature Engineering?

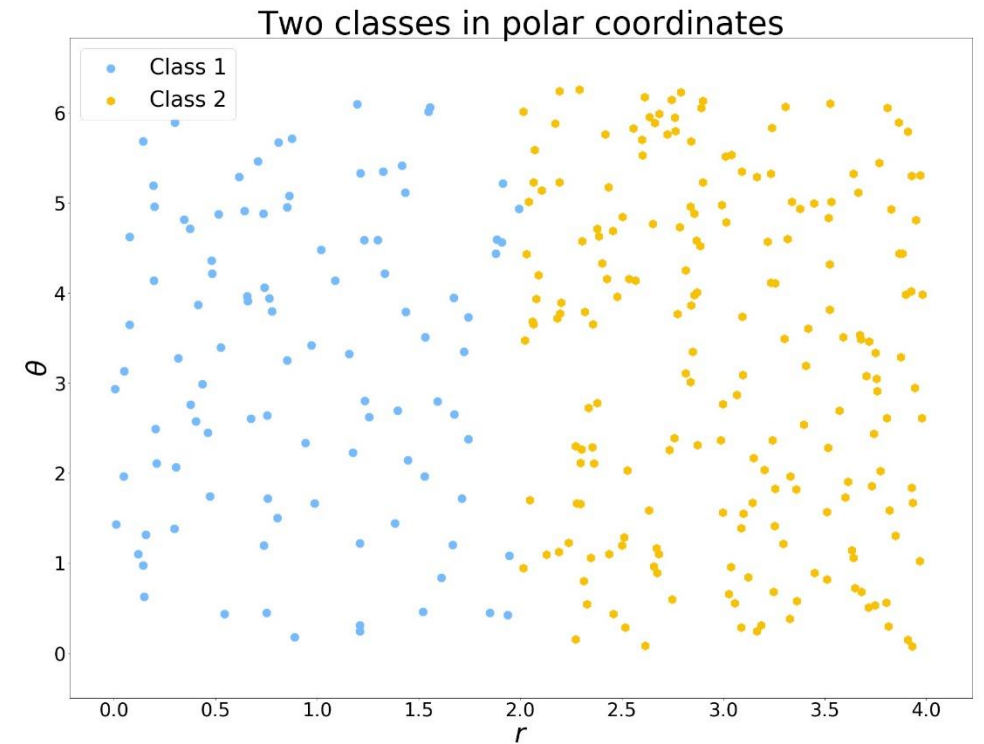
- Even the raw dataset has features. Most of the time, the data will be in the form of a table.
- Each column is a feature. But these features may not produce the best results from the algorithm.
- Modifying, deleting and combining these features results in a new set that is more adept at training the algorithm.
- Feature engineering in machine learning is more than selecting the appropriate features and transforming them.
- Not only does feature engineering prepare the dataset to be compatible with the algorithm, but it also improves the performance of the machine learning models.



## Example for the importance of Feature Engineering



- Original Data: non-linear boundary and not convenient for many machine learning algorithms



- Derived Data: linear boundary and will work with almost all machine learning algorithms





- A common workflow for model development:
  - use pandas for data loading and cleaning.
  - a modeling library to build the model.
- The point of contact between pandas & other analysis libraries:
  - NumPy arrays
- To turn a DataFrame into a NumPy array:
  - Use the .values property:
- To convert back to a DataFrame:
  - pass a two-dimensional ndarray with optional column names

```
df2 = pd.DataFrame(data.values,
                    columns=['one', 'two', 'three'])
df2
```

	one	two	three
0	1.0	0.01	-1.5
1	2.0	-0.01	0.0
2	3.0	0.25	3.6
3	4.0	-4.10	1.3
4	5.0	0.00	-2.0

```
data = pd.DataFrame(
    {'x0': [1, 2, 3, 4, 5],
     'x1': [0.01, -0.01, 0.25, -4.1, 0.],
     'y': [-1.5, 0., 3.6, 1.3, -2.]})
data
```

	x0	x1	y
0	1	0.01	-1.5
1	2	-0.01	0.0
2	3	0.25	3.6
3	4	-4.10	1.3
4	5	0.00	-2.0

```
data.columns
```

```
Index(['x0', 'x1', 'y'], dtype='object')
```

```
data.values
```

```
array([[ 1.  ,  0.01, -1.5 ],
       [ 2.  , -0.01,  0.  ],
       [ 3.  ,  0.25,  3.6 ],
       [ 4.  , -4.1 ,  1.3 ],
       [ 5.  ,  0.  , -2.  ]])
```



## ➤ Note:

- The `.values` attribute is intended to be used when your data is homogeneous - for example, all numeric types.
- If you have heterogeneous data, the result will be an `ndarray` of Python objects.

```
df3 = data.copy()
df3['strings'] = ['a', 'b', 'c', 'd', 'e']
df3
```

	x0	x1	y	strings
0	1	0.01	-1.5	a
1	2	-0.01	0.0	b
2	3	0.25	3.6	c
3	4	-4.10	1.3	d
4	5	0.00	-2.0	e

```
df3.values
```

```
array([[1, 0.01, -1.5, 'a'],
       [2, -0.01, 0.0, 'b'],
       [3, 0.25, 3.6, 'c'],
       [4, -4.1, 1.3, 'd'],
       [5, 0.0, -2.0, 'e']], dtype=object)
```

- If you want to use a subset of the column

```
model_cols = ['x0', 'x1']
data.loc[:, model_cols].values
```

```
array([[ 1. ,  0.01],
       [ 2. , -0.01],
       [ 3. ,  0.25],
       [ 4. , -4.1 ],
       [ 5. ,  0.  ]])
```

```
data[model_cols].values
```

```
array([[ 1. ,  0.01],
       [ 2. , -0.01],
       [ 3. ,  0.25],
       [ 4. , -4.1 ],
       [ 5. ,  0.  ]])
```



```
data['category'] = pd.Categorical(
    ['a', 'b', 'a', 'a', 'b'],
    categories=['a', 'b'])
data
```

	x0	x1	y	category
0	1	0.01	-1.5	a
1	2	-0.01	0.0	b
2	3	0.25	3.6	a
3	4	-4.10	1.3	a
4	5	0.00	-2.0	b

```
dummies = pd.get_dummies(data.category,
                          prefix='category')
dummies
```

	category_a	category_b
0	1	0
1	0	1
2	1	0
3	1	0
4	0	1

## Dealing with Categorical data

- If we wanted to replace the 'category' column with dummy variables:
  - create dummy variables
  - drop the 'category' column
  - join the result

```
data_with_dummies = data.drop('category',
                              axis=1).join(dummies)
data_with_dummies
```

	x0	x1	y	category_a	category_b
0	1	0.01	-1.5	1	0
1	2	-0.01	0.0	0	1
2	3	0.25	3.6	1	0
3	4	-4.10	1.3	1	0
4	5	0.00	-2.0	0	1



MAGIC CODE INSTITUTE

## Section 6.2 | Creating Model Descriptions with Patsy

■



## Patsy:

- A Python library for describing statistical models (especially linear models) with a small string-based “formula syntax”.
- well supported for specifying linear models in `statsmodels`
  - “`y ~ x0 + x1`”
  - `patsy.dmatrices` function takes a formula string along with a dataset (a `DataFrame` or a `dict` of arrays) and produces design matrices for a linear model

```
data = pd.DataFrame({
    'x0': [1, 2, 3, 4, 5],
    'x1': [0.01, -0.01, 0.25, -4.1, 0.],
    'y': [-1.5, 0., 3.6, 1.3, -2.]})
data
```

	x0	x1	y
0	1	0.01	-1.5
1	2	-0.01	0.0
2	3	0.25	3.6
3	4	-4.10	1.3
4	5	0.00	-2.0

```
import patsy
y, X = patsy.dmatrices('y ~ x0 + x1', data)
```

X

DesignMatrix with shape (5, 3)

	Intercept	x0	x1
1	1	0.01	
1	2	-0.01	
1	3	0.25	
1	4	-4.10	
1	5	0.00	

Terms:

'Intercept' (column 0)  
'x0' (column 1)  
'x1' (column 2)

y

DesignMatrix with shape (5, 1)

	y
-1.5	
0.0	
3.6	
1.3	
-2.0	

Terms:

'y' (column 0)



- These Patsy `DesignMatrix` instances are `NumPy ndarray`s with additional metadata

```
np.asarray(X)
```

```
array([[ 1. ,  1. ,  0.01],
       [ 1. ,  2. , -0.01],
       [ 1. ,  3. ,  0.25],
       [ 1. ,  4. , -4.1 ],
       [ 1. ,  5. ,  0.  ]])
```

```
np.asarray(y)
```

```
array([[ -1.5],
       [  0. ],
       [  3.6],
       [  1.3],
       [ -2. ]])
```

```
X.design_info.column_names
```

```
['Intercept', 'x0', 'x1']
```

```
y.design_info.column_names
```

```
['y']
```

- The `Intercept` term is a convention for linear models like **ordinary least squares (OLS) regression**.
- You can suppress the intercept by adding the term `" + 0 "` to the model

```
patsy.dmatrices('y ~ x0 + x1 + 0', data)[1]
# X here
```

DesignMatrix with shape (5, 2)

	x0	x1
1	0.01	
2	-0.01	
3	0.25	
4	-4.10	
5	0.00	

Terms:

'x0' (column 0)

'x1' (column 1)





```
np.linalg.lstsq(X, y)
```

```
(array([[ 0.31290976],
        [-0.07910564],
        [-0.26546384]]),
 array([19.63791494]),
 3,
 array([8.03737688, 3.38335321, 0.90895207]))
```

- The Patsy objects can be passed directly into algorithms

```
coef, resid, _, _ = np.linalg.lstsq(X, y)
```

```
coef
```

```
array([[ 0.31290976],
        [-0.07910564],
        [-0.26546384]])
```

- Get model coefficients

```
coef.squeeze()
```

```
array([ 0.31290976, -0.07910564, -0.26546384])
```

```
coef = pd.Series(coef.squeeze(), index=X.design_info.column_names)
coef
```

- reattach the model column names to the fitted coefficients to obtain a Series

```
Intercept    0.312910
x0           -0.079106
x1           -0.265464
dtype: float64
```





- Data Transformations in Patsy Formulas
  - **Patsy** will try to find the functions you use in the enclosing scope
  
- Some commonly used variable transformations include:
  - Standardizing (to mean 0 and variance 1):
  - Centering (subtracting the mean).

```
y, X = patsy.dmatrices('y ~ x0 + np.log(np.abs(x1) + 1)', data)
X
```

DesignMatrix with shape (5, 3)

Intercept	x0	np.log(np.abs(x1) + 1)
1	1	0.00995
1	2	0.00995
1	3	0.22314
1	4	1.62924
1	5	0.00000

Terms:

'Intercept' (column 0)  
 'x0' (column 1)  
 'np.log(np.abs(x1) + 1)' (column 2)

```
y, X = patsy.dmatrices('y ~ standardize(x0) + center(x1)', data)
X
```

DesignMatrix with shape (5, 3)

Intercept	standardize(x0)	center(x1)
1	-1.41421	0.78
1	-0.70711	0.76
1	0.00000	1.02
1	0.70711	-3.33
1	1.41421	0.77

Terms:

'Intercept' (column 0)  
 'standardize(x0)' (column 1)  
 'center(x1)' (column 2)



➤ Try some other numpy functions

FUNCTION	DESCRIPTION
<b>expm1()</b>	Calculate $\exp(x) - 1$ for all elements in the array.
<b>exp2()</b>	Calculate $2^{**}p$ for all p in the input array.
<b>log10()</b>	Return the base 10 logarithm of the input array, element-wise.
<b>log2()</b>	Base-2 logarithm of x.
<b>log1p()</b>	Return the natural logarithm of one plus the input array, element-wise.
<b>logaddexp()</b>	Logarithm of the sum of exponentiations of the inputs.
<b>logaddexp2()</b>	Logarithm of the sum of exponentiations of the inputs in base-2.

FUNCTION	DESCRIPTION
<b>tan()</b>	Compute tangent element-wise.
<b>arcsin()</b>	Inverse sine, element-wise.
<b>arccos()</b>	Trigonometric inverse cosine, element-wise.
<b>arctan()</b>	Trigonometric inverse tangent, element-wise.
<b>arctan2()</b>	Element-wise arc tangent of x1/x2 choosing the quadrant correctly.
<b>degrees()</b>	Convert angles from radians to degrees.
<b>rad2deg()</b>	Convert angles from radians to degrees.
<b>deg2rad</b>	Convert angles from degrees to radians.
<b>radians()</b>	Convert angles from degrees to radians.
<b>hypot()</b>	Given the “legs” of a right triangle, return its hypotenuse.
<b>unwrap()</b>	Unwrap by changing deltas between values to $2\pi$ complement.



### `patsy.center(x)`

A stateful transform that centers input data, i.e., subtracts the mean.

If input has multiple columns, centers each column separately.

Equivalent to `standardize(x, rescale=False)`

### `patsy.standardize(x, center=True, rescale=True, ddof=0)`

A stateful transform that standardizes input data, i.e. it subtracts the mean and divides by the sample standard deviation.

Either centering or rescaling or both can be disabled by use of keyword arguments. The *ddof* argument controls the delta degrees of freedom when computing the standard deviation (cf. `numpy.std()`). The default of `ddof=0` produces the maximum likelihood estimate; use `ddof=1` if you prefer the square root of the unbiased estimate of the variance.

If input has multiple columns, standardizes each column separately.

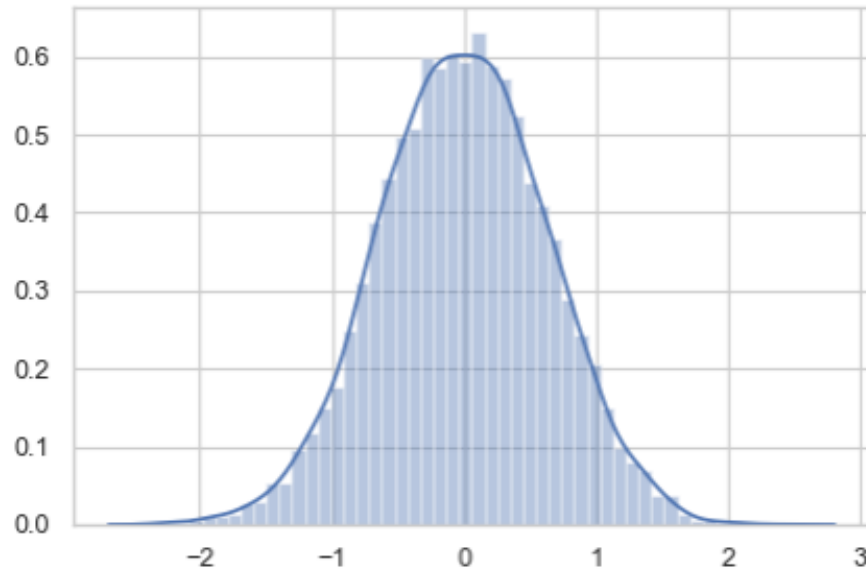




- Import modules
- `dnorm`: helper function to generate random numbers following normal distribution
- Y is constructed as:
  - $Y = X * \text{beta} + \text{eps}$  (noise)

```
fig = plt.figure()
sns.distplot(X[:,0])
```

Figure 1



```
import statsmodels.api as sm
import statsmodels.formula.api as smf
```

```
def dnorm(mean, variance, size=1):
    if isinstance(size, int):
        size = size,
    return mean + np.sqrt(variance) * np.random.randn(*size)
```

```
# For reproducibility
np.random.seed(12345)
```

```
N = 10000
X = np.c_[dnorm(0, 0.4, size=N),
          dnorm(0, 0.6, size=N),
          dnorm(0, 0.2, size=N)]
eps = dnorm(0, 0.1, size=N)
beta = [0.1, 0.3, 0.5]
y = np.dot(X, beta) + eps
```

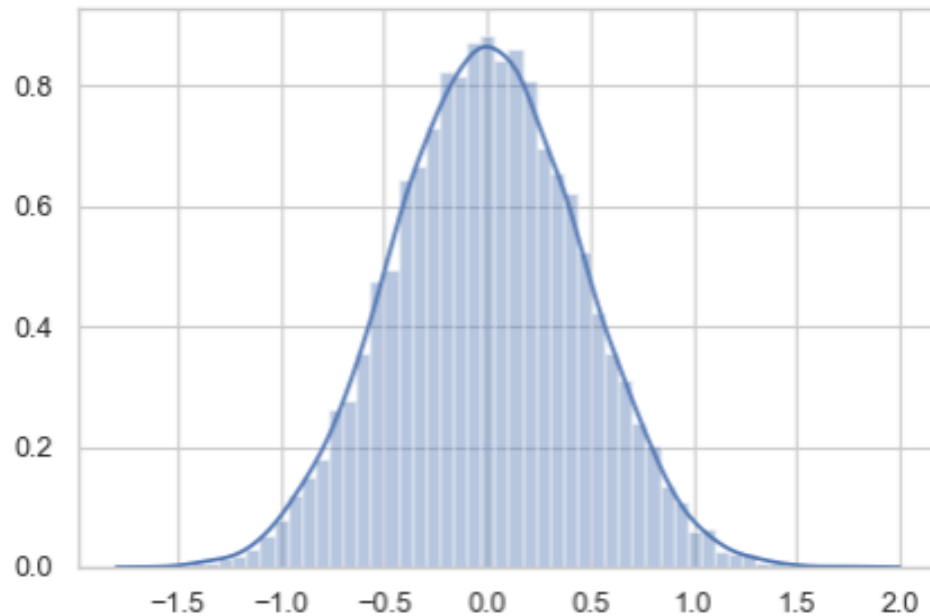
- Normally distributed with:
  - mean = 0
  - Variance = 0.4



- Normally distributed with:
  - mean = 0
  - Variance = 0.6 (0.2)

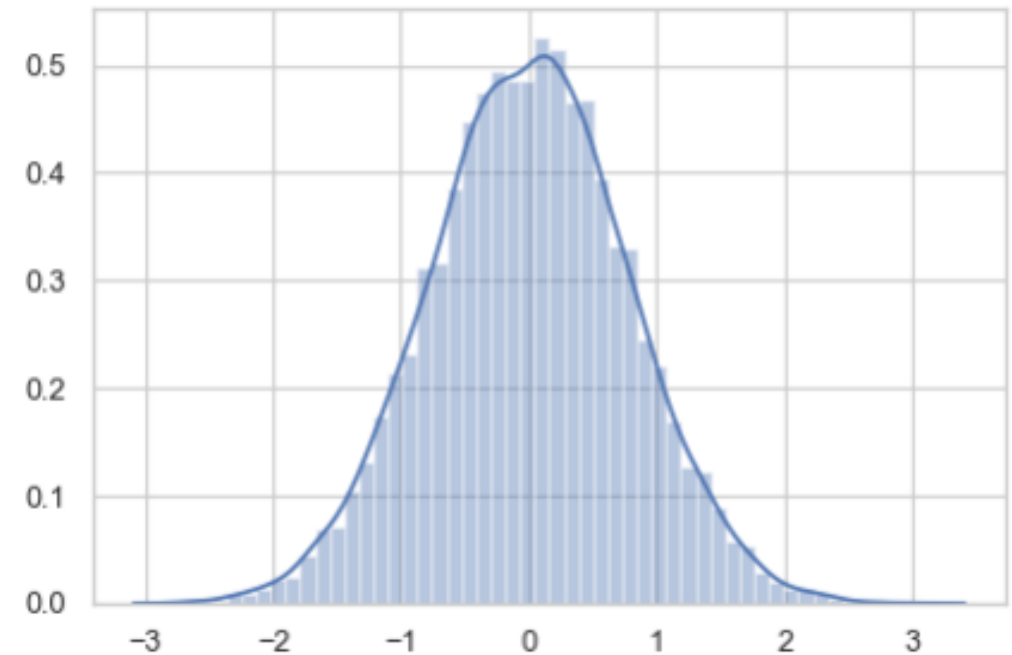
```
fig = plt.figure()  
sns.distplot(X[:,2])
```

Figure 3



```
fig = plt.figure()  
sns.distplot(X[:,1])
```

Figure 2



```
model = sm.OLS(y, X)  
results = model.fit()  
results.params
```

```
array([0.09166321, 0.30220827, 0.50473953])
```





```
print(results.summary())
```

### OLS Regression Results

```
=====
Dep. Variable:          y      R-squared (uncentered):      0.521
Model:                OLS      Adj. R-squared (uncentered):    0.521
Method:             Least Squares      F-statistic:          3629.
Date:               Sat, 04 Apr 2020      Prob (F-statistic):    0.00
Time:               16:40:52      Log-Likelihood:       -2708.3
No. Observations:    10000      AIC:                  5423.
Df Residuals:        9997      BIC:                  5444.
Df Model:             3
Covariance Type:     nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
x1	0.0917	0.005	18.242	0.000	0.082	0.102
x2	0.3022	0.004	73.569	0.000	0.294	0.310
x3	0.5047	0.007	71.712	0.000	0.491	0.519

```
=====
Omnibus:              3.283      Durbin-Watson:          1.979
Prob(Omnibus):         0.194      Jarque-Bera (JB):       3.314
Skew:                 -0.041      Prob(JB):               0.191
Kurtosis:              2.966      Cond. No.                1.71
=====
```





```
results = smf.ols('y ~ col0 + col1 + col2', data=data).fit()
results.params
```

```
Intercept    0.000858
col0         0.091678
col1         0.302207
col2         0.504745
dtype: float64
```

```
data = pd.DataFrame(X, columns=['col0', 'col1', 'col2'])
data['y'] = y
data
```

	col0	col1	col2	y
0	-0.129468	1.493446	-0.258437	0.142901
1	0.302910	-0.895914	0.375710	-0.329110
2	-0.328522	-0.302596	0.313175	0.288018
3	-0.351475	0.310462	0.041303	-0.251139
4	1.243269	-0.677377	0.509666	0.127453
...	...	...	...	...
9995	-0.545716	0.280965	-0.707475	-0.295103
9996	1.361230	2.058392	-0.622691	0.461859
9997	-0.004241	1.140471	-0.331893	0.422961
9998	-0.768258	0.800278	-0.419750	-0.061137
9999	0.414251	2.285421	0.017529	0.491187

10000 rows × 4 columns



```
data = pd.DataFrame(X, columns=['col0', 'col1', 'col2'])
data['y'] = y
data
```

	col0	col1	col2	y
0	-0.129468	1.493446	-0.258437	0.142901
1	0.302910	-0.895914	0.375710	-0.329110
2	-0.328522	-0.302596	0.313175	0.288018
3	-0.351475	0.310462	0.041303	-0.251139
4	1.243269	-0.677377	0.509666	0.127453
...	...	...	...	...
9995	-0.545716	0.280965	-0.707475	-0.295103
9996	1.361230	2.058392	-0.622691	0.461859
9997	-0.004241	1.140471	-0.331893	0.422961
9998	-0.768258	0.800278	-0.419750	-0.061137
9999	0.414251	2.285421	0.017529	0.491187

10000 rows × 4 columns

```
results = smf.ols('y ~ col0 + col1 + col2',
                  data=data).fit()
results.params
```

```
Intercept    0.000858
col0         0.091678
col1         0.302207
col2         0.504745
dtype: float64
```

- Suppose instead that all of the model parameters are in a `DataFrame`
- use the `statsmodels` formula API and `Patsy` formula strings





**MAGIC CODE INSTITUTE**

**THANKS FOR LISTENING!!!**

